

TFM - Kaggle House Prices: Advanced Regression Techniques with caret

04 Creación de modelos predictivos con caret

Juan Carlos Santiago Culebras

2019-09-22

En esta fase aplicaremos distintos algoritmos de machine learning para generar modelos de regresión, que sean capaces de predecir la variable objetivo (SalePrice).

Existen multitud de algoritmos ya implementados para entrenar modelos de regresión, el paquete Caret simplifica la llamada ofreciendo un interfaz único.

Primeros pasos

Librerías

Realizamos la carga de las librerías necesarias

```
if(!is.element("dplyr", installed.packages()[, 1]))
  install.packages("dplyr", repos = 'http://cran.us.r-project.org')
library(dplyr)

if(!is.element("tidyr", installed.packages()[, 1]))
  install.packages("tidyr", repos = 'http://cran.us.r-project.org')
library(tidyr)

if(!is.element("ggplot2", installed.packages()[, 1]))
  install.packages("ggplot2", repos = 'http://cran.us.r-project.org')
library(ggplot2)

if(!is.element("grid", installed.packages()[, 1]))
  install.packages("grid", repos = 'http://cran.us.r-project.org')
library(grid)

if(!is.element("gridExtra", installed.packages()[, 1]))
  install.packages("gridExtra", repos = 'http://cran.us.r-project.org')
library(gridExtra)

if(!is.element("ggpubr", installed.packages()[, 1]))
  install.packages("ggpubr", repos = 'http://cran.us.r-project.org')
library(ggpubr)

if(!is.element("tibble", installed.packages()[, 1]))
  install.packages("tibble", repos = 'http://cran.us.r-project.org')
library(tibble)

if(!is.element("randomForest", installed.packages()[, 1]))
```

```

install.packages("randomForest", repos = 'http://cran.us.r-project.org')
library(randomForest)

if(!is.element("recipes", installed.packages()[, 1]))
  install.packages("recipes", repos = 'http://cran.us.r-project.org')
library(recipes)

if(!is.element("caret", installed.packages()[, 1]))
  install.packages("caret", repos = 'http://cran.us.r-project.org')
library(caret)

if(!is.element("kernlab", installed.packages()[, 1]))
  install.packages("kernlab", repos = 'http://cran.us.r-project.org')
library(kernlab)

if(!is.element("ranger", installed.packages()[, 1]))
  install.packages("ranger", repos = 'http://cran.us.r-project.org')
library(ranger)

if(!is.element("gbm", installed.packages()[, 1]))
  install.packages("gbm", repos = 'http://cran.us.r-project.org')
library(gbm)

if(!is.element("e1071", installed.packages()[, 1]))
  install.packages("e1071", repos = 'http://cran.us.r-project.org')
library(e1071)

if(!is.element("elasticnet", installed.packages()[, 1]))
  install.packages("elasticnet", repos = 'http://cran.us.r-project.org')
library(elasticnet)

if(!is.element("xgboost", installed.packages()[, 1]))
  install.packages("xgboost", repos = 'http://cran.us.r-project.org')
library(xgboost)

if(!is.element("glmnet", installed.packages()[, 1]))
  install.packages("glmnet", repos = 'http://cran.us.r-project.org')
library(glmnet)

```

Funciones

```

fnEstudioModelo <- function ( modelo , estudioParam = TRUE){

  # modelo
  # modelo$finalModel

  p1 <- ggplot(data = modelo$resample, aes(x = RMSE)) +
    geom_density(alpha = 0.5, fill = "gray50") +
    geom_vline(xintercept = mean(modelo$resample$RMSE),
               linetype = "dashed") +
    theme_bw()

```

```

p2 <- ggplot(data = modelo$resample, aes(x = 1, y = RMSE)) +
  geom_boxplot(outlier.shape = NA, alpha = 0.5, fill = "gray50") +
  geom_jitter(width = 0.05) +
  labs(x = "") +
  theme_bw() +
  theme(axis.text.x = element_blank(), axis.ticks.x = element_blank())

#Estudio de hiperparámetros
if (estudioParam){
  p3 <- plot(modelo)
}

# Error de test
predicciones <- predict(modelo
                        , newdata = dsTrain.CV
                        , type = "raw")

# RMSE(predicciones, dsTrain.CV$SalePrice)
# MAE(predicciones, dsTrain.CV$SalePrice)
# R2(predicciones, dsTrain.CV$SalePrice, form = "traditional")

t1 <- capture.output(summary(modelo$resample$RMSE, digits=3))
t1 <- paste("Summary resample$RMSE", " ", paste(t1, collapse="\n"), sep = "\n")
t1 <- text_grob(t1, size = 10)

t2 <- capture.output(postResample(pred = predicciones, obs = dsTrain.CV$SalePrice))
t2 <- paste("Error de test", " ", paste(t2, collapse="\n"), sep = "\n")
t2 <- text_grob(t2, size = 10)

t3 <- capture.output(modelo$finalModel)
t3 <- text_grob(paste(t3, collapse="\n"), size = 9)

grid.arrange(t3, top="Modelo final")
grid.arrange(p1, p2, t1, t2, nrow = 2, top="RMSE obtenido en la validación")

if (estudioParam){
  grid.arrange(p3, nrow = 1, top="Evolución del RMSE del modelo en función de hiperparámetros")
}
}

```

Cargamos datos

Partimos de un dataset, construido en las etapas anteriores, con los datos ya preparados y que contiene solo las variables predictivas.

```

# Conjunto seleccionado en paso anterior

strOrigenF2 <- 'F02_03_dsDataAll_Recipe'
strOrigenF3 <- 'F03_11_dsDataSelVar_rfe_MejorRendimiento_top18'

```

```

file <- paste('./F03_SelPredictores/',strOrigenF2,'/',strOrigenF3,'.RData',sep='')

load(file)

dirSalida <- paste('./F04_Modelos/',strOrigenF2,sep='')

if (!file.exists(dirSalida)){
  dir.create(file.path(dirSalida))
}

dirSalida <- paste('./F04_Modelos/',strOrigenF2,'/',strOrigenF3,sep='')

if (!file.exists(dirSalida)){
  dir.create(file.path(dirSalida))
}

rm(file)

```

Lectura de modelos ya entrenados si se realiza es estudio posteriormente

```

# dir <- './F04_Modelos/F02_03_dsDataAll_Recipe/F03_15_dsDataSelVar_Completo/'
# load(paste(dir, 'modelo_gbm.RData', sep=''))
# load(paste(dir, 'modelo_glm.RData', sep=''))
# load(paste(dir, 'modelo_glmnet.RData', sep=''))
# load(paste(dir, 'modelo_Knn.RData', sep=''))
# load(paste(dir, 'modelo_lasso.RData', sep=''))
# load(paste(dir, 'modelo_lm.RData', sep=''))
# load(paste(dir, 'modelo_rf.RData', sep=''))
# load(paste(dir, 'modelo_sumlineal.RData', sep=''))
# load(paste(dir, 'modelo_sumRadial.RData', sep=''))
# load(paste(dir, 'modelo_XGBoost.RData', sep=''))

```

Separamos los datos

Optenemos 4 dataset

dsTrain - Que a su vez se divide en dsTrain.training dsTrain.CV

dsTest

```

dsTrain <- dsDataAllVarSel %>%
  filter(indTrain == 1) %>%
  select(SalePrice, everything()) %>%
  select(-c(Id,indTrain))

dim(dsTrain)

```

```
## [1] 1458 19
```

```

set.seed(123)
iTrain <- createDataPartition(y=dsTrain$SalePrice, p=0.7, list=F)

```

```
dsTrain.training <- dsTrain[iTrain, ]
dsTrain.CV       <- dsTrain[-iTrain, ]

dsTest <- dsDataAllVarSel %>%
  filter(indTrain == 0) %>%
  select(SalePrice, everything())
```

Modelos

En esta sección se entrenarán distintos modelos para evaluar cual puede ser el mejor.

Sobre cada modelo se realizará:

- Entrenamiento
- Ajuste de hiperparámetros
- Evaluación mediante validación cruzada

Definimos tipo de entrenamiento

```
particiones <- 5
repeticiones <- 5

# Entrenamiento con conjunto de hiperparametros
fitControl <- trainControl(method = "repeatedcv",
                           number = particiones,
                           repeats = repeticiones,
                           returnResamp = "final",
                           verboseIter = FALSE)
```

Regresión lineal

En estos modelos se busca una función con los predictores como variables y una combinación de pesos que multiplicados por las variables den como resultado un modelo para la variable objetivo.

Estos algoritmos son muy rápidos y responden bien cuando el número de predictores es alto.

En nuestro caso hemos seleccionado dos ejemplos

Linear Regression

Regresión lineal, este modelo es el más sencillo de probar y me ha servido como línea base para ir evaluando el resto de los modelos

```
t <- proc.time() # Inicia el cronómetro
modelo_lm <- train(SalePrice ~ .
                  , data = dsTrain.training
                  , method = "lm"
                  #, tuneGrid = hiperparametros
                  , tuneLength = 10
                  , metric = "RMSE"
                  , trControl = fitControl)
proc.time()-t    # Detiene el cronómetro
```

```
## user system elapsed
## 0.95 0.02 0.97
```

```
# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_lm', '.RData', sep='')
save(modelo_lm, file = fileOuput)

modelo_lm
```

```
## Linear Regression
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 818, 820, 817, 819, 818, 819, ...
## Resampling results:
##
## RMSE      Rsquared   MAE
## 0.1299566  0.8976907  0.0931044
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
# Presento estudio
fnEstudioModelo(modelo_lm, estudioParam = FALSE)
```



Modelo final

Call:
lm(formula = .outcome ~ ., data = dat)

Coefficients:

(Intercept)	GrLivArea	OverallQual	TotalBsmtSF	X1stFlrSF
12.0216986	0.1191885	0.0953092	0.0484734	0.0036513
BsmtFinSF1	LotArea	GarageArea	X2ndFlrSF	OverallCond
0.0326487	0.0367879	0.0168238	0.0065776	0.0578678
FireplaceQu	YearBuilt	KitchenQual	YearRemodAdd	GarageCars
0.0243120	0.0822681	0.0190883	0.0255445	0.0323463
ExterQual	BsmtFinType1	TotRmsAbvGrd	Fireplaces	
0.0121054	0.0093991	-0.0008169	0.0119496	

RMSE obtenido en la validación



Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.114	0.125	0.131	0.130	0.136	0.147

Error de test

RMSE	Rsquared	MAE
0.12362414	0.89944770	0.08982126

Generalized Linear Model

```
hiperparametros <- data.frame(parameter = "none")

t <- proc.time() # Inicia el cronómetro
modelo_glm <- train(SalePrice ~ .
                    , data = dsTrain.training
                    , method = "glm"
                    , tuneGrid = hiperparametros
                    , metric = "RMSE"
                    , trControl = fitControl)
proc.time()-t    # Detiene el cronómetro

##      user  system elapsed
##      0.99    0.03    1.01

# Guardo resultado del calculo
fileOuput <- paste(dirSalida, '/', 'modelo_glm', '.RData', sep='')
save(modelo_glm, file = fileOuput)

modelo_glm
```

```
## Generalized Linear Model
```

```
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 819, 818, 819, 819, 817, 819, ...
## Resampling results:
##
## RMSE      Rsquared  MAE
## 0.1297683  0.898214  0.09318374
```

```
# Presento estudio
fnEstudioModelo(modelo_glm, estudioParam = FALSE)
```

Modelo final

Call: NULL


Coefficients:

(Intercept)	GrLivArea	OverallQual	TotalBsmtSF	X1stFlrSF
12.0216986	0.1191885	0.0953092	0.0484734	0.0036513
BsmtFinSF1	LotArea	GarageArea	X2ndFlrSF	OverallCond
0.0326487	0.0367879	0.0168238	0.0065776	0.0578678
FireplaceQu	YearBuilt	KitchenQual	YearRemodAdd	GarageCars
0.0243120	0.0822681	0.0190883	0.0255445	0.0323463
ExterQual	BsmtFinType1	TotRmsAbvGrd	Fireplaces	
0.0121054	0.0093991	-0.0008169	0.0119496	

Degrees of Freedom: 1022 Total (i.e. Null); 1004 Residual

Null Deviance: 167.8

Residual Deviance: 16.62 AIC: -1271



Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.107	0.116	0.132	0.130	0.141	0.159

Error de test

RMSE	Rsquared	MAE
0.12362414	0.89944770	0.08982126

Support Vector Machines

Aunque Las máquinas de vectores soporte fueron pensadas para resolver problemas de clasificación también pueden adaptarse para resolver problemas de regresión, estos modelos dan bastante buenos resultados cuando la variable objetivo no es separables linealmente dentro del espacio vectorial de los predictores y evitan en gran medida el problema del sobreajuste a los ejemplos de entrenamiento, por ello es una buena elección para este problema.

Las máquinas de soporte utilizan una función denominada Kernel para la búsqueda del hiperplano de separación, para ello mapean los datos en espacios de dimensiones superiores con la esperanza de que en este espacio de dimensiones superiores los datos puedan separarse más fácilmente o estar mejor estructurados.

En nuestro caso hemos probado con dos modelos con funciones Kernel distintas:

Support Vector Machines with Linear Kernel

Permite solo seleccionar líneas (o hiperplanos)

```
hiperparametros <- data.frame(C = c(0.0001, 0.001, 0.01, 0.1, 0.5))

t <- proc.time() # Inicia el cronómetro
modelo_svmlineal <- train(SalePrice ~ .
  , data = dsTrain.training
  , method = "svmLinear"
  , tuneGrid = hiperparametros)
```

```

, metric = "RMSE"
, trControl = fitControl
, scale = FALSE )
proc.time()-t    # Detiene el cronómetro

```

```

##    user  system elapsed
##    9.30    0.11    9.41

```

```

# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_svmlineal', '.RData', sep='')
save(modelo_svmlineal, file = fileOuput)

modelo_svmlineal

```

```

## Support Vector Machines with Linear Kernel
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 819, 818, 819, 819, 817, 819, ...
## Resampling results across tuning parameters:
##
##  C      RMSE      Rsquared  MAE
##  1e-04  0.1954332  0.8753939  0.14056770
##  1e-03  0.1343825  0.8933460  0.09532633
##  1e-02  0.1303350  0.8974270  0.09287770
##  1e-01  0.1304045  0.8972558  0.09319034
##  5e-01  0.1304532  0.8971945  0.09324501
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was C = 0.01.

```

```

# Presento estudio
fnEstudioModelo(modelo_svmlineal)

```

Modelo final

Support Vector Machine object of class "ksvm"



SV type: eps-svr (regression)
parameter : epsilon = 0.1 cost C = 0.01

Linear (vanilla) kernel function.

Number of Support Vectors : 349

Objective Function Value : -0.2898
Training error : 0.016503

RMSE obtenido en la validación




Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.112	0.124	0.130	0.130	0.137	0.149

Error de test

RMSE	Rsquared	MAE
0.12266020	0.90033055	0.08949161

Evolución del RMSE del modelo en función de hiperparámetros



Support Vector Machines with Radial Basis Function Kernel

Permiten seleccionar círculos (o hipersferas)

```
hiperparametros <- expand.grid(sigma = c(0.0005, 0.001, 0.005)
                                ,C = c(1, 20, 50, 100, 150, 200))
```

```
t <- proc.time() # Inicia el cronómetro
modelo_svmRadial <- train(SalePrice ~ .
                          , data = dsTrain.training
                          , method = "svmRadial"
                          , tuneGrid = hiperparametros
                          , metric = "RMSE"
                          , trControl = fitControl)
proc.time()-t    # Detiene el cronómetro
```

```
##    user  system elapsed
##   94.31    1.00   95.32
```

```
# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_svmRadial', '.RData', sep='')
save(modelo_svmRadial, file = fileOuput)
```

```
modelo_svmRadial
```

```

## Support Vector Machines with Radial Basis Function Kernel
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 818, 819, 819, 819, 817, 820, ...
## Resampling results across tuning parameters:
##
##  sigma  C    RMSE      Rsquared  MAE
##  5e-04   1  0.1412523  0.8896820  0.09829255
##  5e-04  20  0.1290817  0.9001458  0.08982726
##  5e-04  50  0.1276003  0.9022077  0.08862283
##  5e-04 100  0.1266084  0.9036535  0.08779753
##  5e-04 150  0.1263153  0.9041051  0.08744767
##  5e-04 200  0.1260968  0.9043727  0.08713926
##  1e-03   1  0.1355441  0.8939751  0.09411630
##  1e-03  20  0.1273470  0.9027064  0.08835688
##  1e-03  50  0.1262710  0.9042150  0.08732574
##  1e-03 100  0.1263330  0.9038234  0.08704188
##  1e-03 150  0.1266452  0.9032326  0.08724417
##  1e-03 200  0.1268319  0.9028378  0.08738384
##  5e-03   1  0.1319973  0.8974042  0.09133254
##  5e-03  20  0.1263575  0.9035468  0.08733542
##  5e-03  50  0.1283594  0.9003121  0.08908754
##  5e-03 100  0.1304894  0.8968835  0.09104080
##  5e-03 150  0.1321543  0.8942250  0.09218969
##  5e-03 200  0.1336779  0.8918734  0.09315568
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 5e-04 and C = 200.

```

```

# Presento estudio
fnEstudioModelo(modelo_svmRadial)

```

Modelo final

Support Vector Machine object of class "ksvm"



SV type: eps-svr (regression)
parameter : epsilon = 0.1 cost C = 200

Gaussian Radial Basis kernel function.
Hyperparameter : sigma = 5e-04

Number of Support Vectors : 673

Objective Function Value : -25485.12
Training error : 0.088932

RMSE obtenido en la validación




Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0944	0.1190	0.1260	0.1260	0.1330	0.1470

Error de test

RMSE	Rsquared	MAE
0.11829999	0.90728931	0.08208742

Evolución del RMSE del modelo en función de hiperparámetros



Arboles de decisión

Estos modelos generan un conjunto de reglas para segmentar el espacio predictor en una serie de regiones simples, generando un árbol de decisiones.

El método es simple y puede servir bien para la interpretación de los datos, pero no tiene una gran precisión en la predicción. Sin embargo, la combinación de una gran cantidad de árboles puede mejorar mucho la predicción.

He realizado pruebas con:

XGBoost

XGBoost ha sido uno de los modelos más utilizados, esto es así porque se adapta fácilmente ya que es muy flexible, se puede usar tanto en regresión como en clasificación. Utiliza una combinación de modelos más simples (árboles de decisión) y potencia los resultados.

La gran desventaja de este modelo es el ajuste de su gran cantidad de parámetros.

```
hiperparametros <- expand.grid(  
  nrounds = seq(from = 200, to = 500, by = 50),  
  eta = c(0.025, 0.05, 0.1, 0.3),  
  max_depth = c(2, 3, 4, 5, 6),  
  gamma = 0,  
  colsample_bytree = 1,  
  min_child_weight = 1,
```

```

subsample = 1
)

t <- proc.time() # Inicia el cronómetro
modelo_XGBoost <- train(SalePrice ~ .
                        , data = dsTrain.training
                        , method = "xgbTree"
                        , tuneGrid = hiperparametros
                        , metric = "RMSE"
                        , trControl = fitControl)
proc.time()-t      # Detiene el cronómetro

##      user  system elapsed
## 1755.94   943.88  1159.72

# Guardo resultado del calculo
fileOuput <- paste(dirSalida, '/', 'modelo_XGBoost', '.RData', sep='')
save(modelo_XGBoost, file = fileOuput)

modelo_XGBoost

```

```

## eXtreme Gradient Boosting
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 818, 818, 819, 818, 819, 818, ...
## Resampling results across tuning parameters:
##
##  eta    max_depth  nrounds  RMSE      Rsquared  MAE
##  0.025   2          200      0.1686581  0.8731809  0.12881586
##  0.025   2          250      0.1451131  0.8807544  0.10432193
##  0.025   2          300      0.1395812  0.8852033  0.09924938
##  0.025   2          350      0.1370025  0.8881210  0.09724826
##  0.025   2          400      0.1353868  0.8901658  0.09609430
##  0.025   2          450      0.1342905  0.8916538  0.09529405
##  0.025   2          500      0.1334867  0.8928052  0.09464642
##  0.025   3          200      0.1613042  0.8821185  0.12312338
##  0.025   3          250      0.1393405  0.8871948  0.09986962
##  0.025   3          300      0.1353842  0.8902102  0.09609558
##  0.025   3          350      0.1339715  0.8919652  0.09478230
##  0.025   3          400      0.1332530  0.8929649  0.09402824
##  0.025   3          450      0.1329221  0.8934370  0.09362958
##  0.025   3          500      0.1327867  0.8936088  0.09339997
##  0.025   4          200      0.1601743  0.8835690  0.12215974
##  0.025   4          250      0.1380655  0.8887987  0.09896259
##  0.025   4          300      0.1345137  0.8913739  0.09514136
##  0.025   4          350      0.1335094  0.8925757  0.09399368
##  0.025   4          400      0.1330467  0.8932232  0.09343265
##  0.025   4          450      0.1328754  0.8934799  0.09315017
##  0.025   4          500      0.1328169  0.8935741  0.09298833

```

##	0.025	5	200	0.1591903	0.8857200	0.12146174
##	0.025	5	250	0.1368515	0.8907843	0.09794958
##	0.025	5	300	0.1335236	0.8928388	0.09422828
##	0.025	5	350	0.1328414	0.8935663	0.09332368
##	0.025	5	400	0.1326986	0.8937403	0.09300233
##	0.025	5	450	0.1326413	0.8938096	0.09292492
##	0.025	5	500	0.1326812	0.8937195	0.09294968
##	0.025	6	200	0.1594463	0.8865336	0.12182967
##	0.025	6	250	0.1369773	0.8910257	0.09783753
##	0.025	6	300	0.1336904	0.8928008	0.09430574
##	0.025	6	350	0.1331369	0.8932146	0.09361945
##	0.025	6	400	0.1330933	0.8931883	0.09343763
##	0.025	6	450	0.1330943	0.8931492	0.09340458
##	0.025	6	500	0.1331904	0.8929761	0.09347716
##	0.050	2	200	0.1354850	0.8898976	0.09620451
##	0.050	2	250	0.1336775	0.8924948	0.09472775
##	0.050	2	300	0.1327201	0.8939320	0.09381334
##	0.050	2	350	0.1325136	0.8942366	0.09341644
##	0.050	2	400	0.1324710	0.8942647	0.09319621
##	0.050	2	450	0.1325616	0.8940889	0.09311694
##	0.050	2	500	0.1325874	0.8940315	0.09296764
##	0.050	3	200	0.1334023	0.8926518	0.09416905
##	0.050	3	250	0.1328463	0.8934245	0.09350269
##	0.050	3	300	0.1327156	0.8935792	0.09320186
##	0.050	3	350	0.1326048	0.8937054	0.09300393
##	0.050	3	400	0.1326765	0.8935638	0.09288112
##	0.050	3	450	0.1328586	0.8932752	0.09293985
##	0.050	3	500	0.1329866	0.8930573	0.09294473
##	0.050	4	200	0.1334315	0.8925425	0.09357676
##	0.050	4	250	0.1330086	0.8931625	0.09306289
##	0.050	4	300	0.1329777	0.8931748	0.09286633
##	0.050	4	350	0.1331852	0.8928528	0.09284632
##	0.050	4	400	0.1333514	0.8925785	0.09293910
##	0.050	4	450	0.1335654	0.8922359	0.09315131
##	0.050	4	500	0.1338918	0.8916876	0.09349162
##	0.050	5	200	0.1336103	0.8923508	0.09358473
##	0.050	5	250	0.1336570	0.8922359	0.09345034
##	0.050	5	300	0.1338288	0.8919428	0.09350994
##	0.050	5	350	0.1340645	0.8915647	0.09371506
##	0.050	5	400	0.1343026	0.8911517	0.09388282
##	0.050	5	450	0.1345193	0.8908108	0.09410619
##	0.050	5	500	0.1346721	0.8905696	0.09427274
##	0.050	6	200	0.1335153	0.8926352	0.09357134
##	0.050	6	250	0.1337190	0.8922427	0.09368625
##	0.050	6	300	0.1340166	0.8917469	0.09399539
##	0.050	6	350	0.1343453	0.8912116	0.09437216
##	0.050	6	400	0.1345510	0.8908722	0.09466792
##	0.050	6	450	0.1347398	0.8905625	0.09488839
##	0.050	6	500	0.1348982	0.8903042	0.09511398
##	0.100	2	200	0.1333515	0.8927973	0.09397286
##	0.100	2	250	0.1332775	0.8928779	0.09372675
##	0.100	2	300	0.1334075	0.8926425	0.09360503
##	0.100	2	350	0.1336216	0.8923290	0.09360180
##	0.100	2	400	0.1338724	0.8919244	0.09360336


##	0.100	2	450	0.1340846	0.8916140	0.09367040
##	0.100	2	500	0.1343894	0.8911481	0.09380467
##	0.100	3	200	0.1338663	0.8916217	0.09329203
##	0.100	3	250	0.1340879	0.8912279	0.09345280
##	0.100	3	300	0.1344295	0.8906736	0.09368090
##	0.100	3	350	0.1348062	0.8901142	0.09384833
##	0.100	3	400	0.1352184	0.8894857	0.09412066
##	0.100	3	450	0.1355371	0.8889974	0.09437732
##	0.100	3	500	0.1358922	0.8884171	0.09467968
##	0.100	4	200	0.1341348	0.8911749	0.09347594
##	0.100	4	250	0.1344864	0.8906205	0.09366788
##	0.100	4	300	0.1347273	0.8902667	0.09397089
##	0.100	4	350	0.1349954	0.8898512	0.09426140
##	0.100	4	400	0.1352229	0.8895045	0.09449260
##	0.100	4	450	0.1354278	0.8891728	0.09471466
##	0.100	4	500	0.1356322	0.8888469	0.09495706
##	0.100	5	200	0.1356595	0.8887608	0.09480588
##	0.100	5	250	0.1360306	0.8881489	0.09522613
##	0.100	5	300	0.1363148	0.8876840	0.09556350
##	0.100	5	350	0.1365382	0.8873184	0.09581678
##	0.100	5	400	0.1366599	0.8871156	0.09598937
##	0.100	5	450	0.1367689	0.8869338	0.09612521
##	0.100	5	500	0.1368374	0.8868169	0.09622034
##	0.100	6	200	0.1359115	0.8885698	0.09561918
##	0.100	6	250	0.1361189	0.8882096	0.09583905
##	0.100	6	300	0.1362945	0.8879167	0.09601290
##	0.100	6	350	0.1364051	0.8877225	0.09613207
##	0.100	6	400	0.1364642	0.8876270	0.09620385
##	0.100	6	450	0.1365035	0.8875583	0.09625234
##	0.100	6	500	0.1365310	0.8875112	0.09629524
##	0.300	2	200	0.1403653	0.8811588	0.09920321
##	0.300	2	250	0.1412117	0.8797652	0.09956681
##	0.300	2	300	0.1423277	0.8779383	0.10031554
##	0.300	2	350	0.1428099	0.8771796	0.10072288
##	0.300	2	400	0.1431633	0.8765749	0.10096368
##	0.300	2	450	0.1435368	0.8759534	0.10125258
##	0.300	2	500	0.1440297	0.8751603	0.10161427
##	0.300	3	200	0.1425781	0.8773955	0.10043838
##	0.300	3	250	0.1431561	0.8764676	0.10100270
##	0.300	3	300	0.1435802	0.8757788	0.10138578
##	0.300	3	350	0.1437411	0.8755398	0.10158734
##	0.300	3	400	0.1439429	0.8752283	0.10187508
##	0.300	3	450	0.1441049	0.8749536	0.10201409
##	0.300	3	500	0.1442558	0.8746984	0.10216487
##	0.300	4	200	0.1439998	0.8751788	0.10241325
##	0.300	4	250	0.1442434	0.8747702	0.10264236
##	0.300	4	300	0.1443293	0.8746230	0.10278538
##	0.300	4	350	0.1443975	0.8745078	0.10285847
##	0.300	4	400	0.1444348	0.8744456	0.10290002
##	0.300	4	450	0.1444367	0.8744466	0.10290763
##	0.300	4	500	0.1444377	0.8744452	0.10290819
##	0.300	5	200	0.1466048	0.8703508	0.10389401
##	0.300	5	250	0.1466497	0.8702677	0.10395129
##	0.300	5	300	0.1466487	0.8702692	0.10395202

```
## 0.300 5 350 0.1466487 0.8702692 0.10395202
## 0.300 5 400 0.1466487 0.8702692 0.10395202
## 0.300 5 450 0.1466487 0.8702692 0.10395202
## 0.300 5 500 0.1466487 0.8702692 0.10395202
## 0.300 6 200 0.1460912 0.8712796 0.10394517
## 0.300 6 250 0.1460912 0.8712796 0.10394517
## 0.300 6 300 0.1460912 0.8712796 0.10394517
## 0.300 6 350 0.1460912 0.8712796 0.10394517
## 0.300 6 400 0.1460912 0.8712796 0.10394517
## 0.300 6 450 0.1460912 0.8712796 0.10394517
## 0.300 6 500 0.1460912 0.8712796 0.10394517
##
## Tuning parameter 'gamma' was held constant at a value of 0
## 1
## Tuning parameter 'min_child_weight' was held constant at a value of
## 1
## Tuning parameter 'subsample' was held constant at a value of 1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nrounds = 400, max_depth = 2,
## eta = 0.05, gamma = 0, colsample_bytree = 1, min_child_weight = 1
## and subsample = 1.
```

```
# Presento estudio
fnEstudioModelo(modelo_XGBoost, estudioParam = FALSE)
```

```
## Modelo final
raw: 153.4 Kb
call:
xgboost::xgb.train(params = list(eta = param$eta, max_depth = param$max_depth,
  gamma = param$gamma, colsample_bytree = param$colsample_bytree,
  min_child_weight = param$min_child_weight, subsample = param$subsample),
  data = x, nrounds = param$nrounds, objective = "reg:linear")
params (as set within xgb.train):
max_depth = "2", gamma = "0", colsample_bytree = "1", min_child_weight = "1", subsample = "1", objective = "reg:line
xgb.attributes:
  niter
  callbacks:
    cb.print.evaluation(period = print_every_n)
  # of features: 18
  niter: 400
  nfeatures : 18
3smtFinSF1 LotArea GarageArea X2ndFlrSF OverallCond FireplaceQu YearBuilt KitchenQual YearRemodAdd Garage
problemType : Regression
tuneValue :
nrounds max_depth eta gamma colsample_bytree min_child_weight
  40 400 2 0.05 0 1 1
  subsample
    40 1
obsLevels : NA
param :
```

RMSE obtenido en la validación



Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.112	0.127	0.131	0.132	0.139	0.151

Error de test

RMSE	Rsquared	MAE
0.12939657	0.89027801	0.09035004

Random Forest

Utiliza una combinación de árboles, en este caso cada árbol depende de los valores de un vector aleatorio.

La ventaja de este método frente a XGBoost es que es más fácil de ajustar, aunque parece menos flexible. También en este modelo se ha detectado un sobreajuste al conjunto de entrenamiento, dando valores bastante buenos en los entrenamientos, pero bastante más altos en test.

```
t <- proc.time() # Inicia el cronómetro
modelo_rf <- train(SalePrice ~ .
                    , data = dsTrain.training
                    , method = "ranger"
                    #, tuneGrid = hiperparametros
                    , tuneLength = 10
                    , metric = "RMSE"
                    , trControl = fitControl
                    , num.trees = 500)
proc.time()-t    # Detiene el cronómetro

##      user  system elapsed
## 1430.61   12.83   225.70

# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_rf', '.RData', sep='')
save(modelo_rf, file = fileOuput)
```

```
modelo_rf
```

```
## Random Forest
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 818, 819, 818, 818, 819, 819, ...
## Resampling results across tuning parameters:
##
##  mtry  splitrule  RMSE      Rsquared  MAE
##  2      variance  0.1404124  0.8890309  0.09656490
##  2      extratrees 0.1457740  0.8840146  0.10111345
##  3      variance  0.1386272  0.8899519  0.09535901
##  3      extratrees 0.1424959  0.8859265  0.09910765
##  5      variance  0.1379895  0.8893221  0.09479262
##  5      extratrees 0.1411151  0.8857655  0.09853838
##  7      variance  0.1383740  0.8880158  0.09508236
##  7      extratrees 0.1404508  0.8857805  0.09823688
##  9      variance  0.1392612  0.8861204  0.09590078
##  9      extratrees 0.1401361  0.8857055  0.09812964
## 10      variance  0.1391568  0.8860387  0.09582295
## 10      extratrees 0.1397212  0.8861267  0.09789956
## 12      variance  0.1403616  0.8837282  0.09669394
## 12      extratrees 0.1398033  0.8858183  0.09812776
## 14      variance  0.1413949  0.8815332  0.09735560
## 14      extratrees 0.1399545  0.8852515  0.09805367
## 16      variance  0.1424339  0.8794758  0.09809577
## 16      extratrees 0.1396919  0.8855031  0.09792580
## 18      variance  0.1443763  0.8756437  0.09931080
## 18      extratrees 0.1399648  0.8849904  0.09807894
##
## Tuning parameter 'min.node.size' was held constant at a value of 5
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were mtry = 5, splitrule =
##  variance and min.node.size = 5.
```

```
# Presento estudio
fnEstudioModelo(modelo_rf)
```

Modelo final


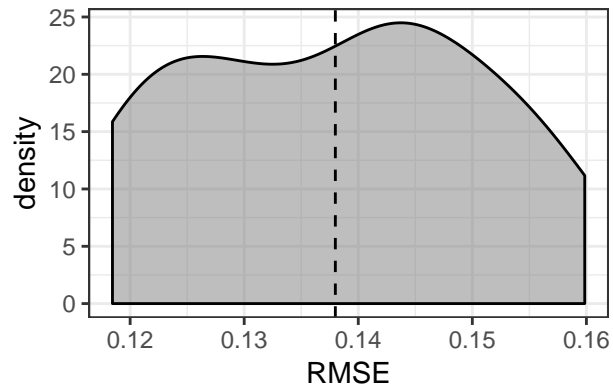
Ranger result

Call:

```
l = x, mtry = min(param$mtry, ncol(x)), min.node.size = param$min.node.size, splitrule = as.character(param$sp
```

Type:	Regression
Number of trees:	500
Sample size:	1023
Number of independent variables:	18
Mtry:	5
Target node size:	5
Variable importance mode:	none
Splitrule:	variance
OOB prediction error (MSE):	0.0185743
R squared (OOB):	0.8868994

RMSE obtenido en la validación




Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.118	0.129	0.140	0.138	0.146	0.160

Error de test

RMSE	Rsquared	MAE
0.13088786	0.88932557	0.09041785

Evolución del RMSE del modelo en función de hiperparámetros



Stochastic Gradient Boosting

GBM realiza un proceso iterativo donde se introducen nuevos modelos que se basan en los errores de las iteraciones anteriores para minimizar el error (aumento de gradiente) de una función objetivo.

Este método es muy versátil, pudiendo resolver una gran variedad de problemas, sus desventajas son que es sensible al sobreajuste, tiene un gran número de hiperparámetros, por lo que es complicado de ajustar y el tiempo de entrenamiento es bastante alto.

```
t <- proc.time() # Inicia el cronómetro
hiperparametros <- expand.grid(interaction.depth = c(2,3),
                              n.trees = c(2000, 3000, 4000),
                              shrinkage = c( 0.01, 0.1),
                              n.minobsinnode = c(2, 5, 10))

t <- proc.time() # Inicia el cronómetro
modelo_gbm <- train(SalePrice ~ .
                    , data = dsTrain.training
                    , method = "gbm"
                    , tuneGrid = hiperparametros
                    , metric = "RMSE"
                    #, distribution
                    , trControl = fitControl
                    , verbose = FALSE # Para que no se muestre cada iteración por pantalla
                    )
proc.time()-t # Detiene el cronómetro
```

```
##      user  system elapsed
## 745.45    0.00  745.64



# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_gbm', '.RData', sep='')
save(modelo_gbm, file = fileOuput)

# Presento estudio
fnEstudioModelo(modelo_gbm)
```

Modelo final

A gradient boosted model with gaussian loss function.
2000 iterations were performed.
There were 18 predictors of which 18 had non-zero influence.

RMSE obtenido en la validación



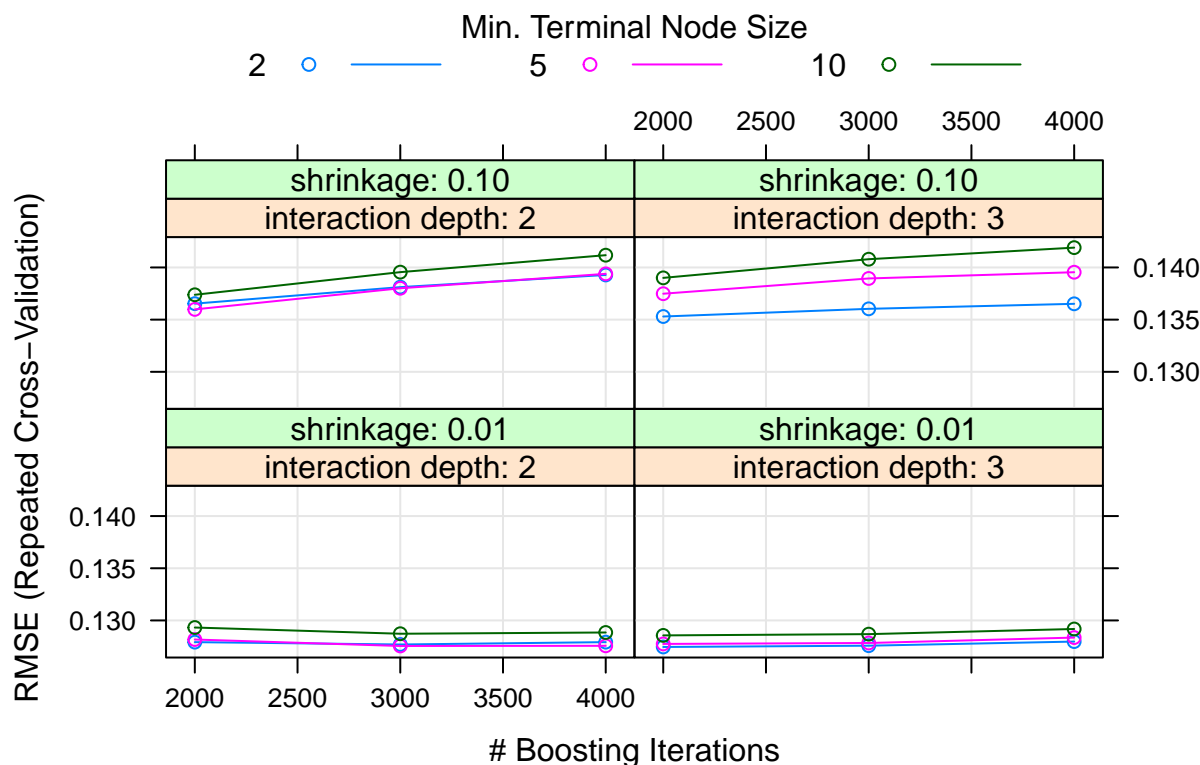
Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.108	0.118	0.130	0.127	0.134	0.152

Error de test

RMSE	Rsquared	MAE
0.12858819	0.89263459	0.08885095

Evolución del RMSE del modelo en función de hiperparámetros



k-Nearest Neighbors

Es un método tanto de clasificación como de regresión, bastante sencillo y supervisado, una característica principal es que está basado en instancia, esto quiere decir que no se genera un modelo real, sino que se guardan las observaciones.

El algoritmo busca las observaciones más cercanas a la que se está tratando y predice el valor de interés mediante los datos que le rodean. El parámetro k indica cuantos puntos “vecinos” se deben de tener en cuenta para ajustar.

KNN tiende a funcionar mejor con dataset pequeños y con pocos predictores, ya que utiliza todo el conjunto de datos para entrenar. Además, es muy costoso tanto en uso de CPU como en memoria.

```
hiperparametros <- data.frame(k = c(3:20))

t <- proc.time() # Inicia el cronómetro
modelo_knn <- train(SalePrice ~ .
  , data = dsTrain.training
  , method = "knn"
  , tuneGrid = hiperparametros
  , metric = "RMSE"
  , trControl = fitControl)
proc.time()-t # Detiene el cronómetro
```

```
## user system elapsed
## 4.57 0.00 4.58
```

```
# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_knn', '.RData', sep='')
save(modelo_knn, file = fileOuput)
```

```
modelo_knn
```



```
## k-Nearest Neighbors
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 819, 818, 819, 819, 817, 817, ...
## Resampling results across tuning parameters:
##
## k RMSE Rsquared MAE
## 3 0.1569542 0.8507359 0.1132504
## 4 0.1540684 0.8573454 0.1114176
## 5 0.1523508 0.8620346 0.1101791
## 6 0.1515969 0.8647515 0.1098665
## 7 0.1514066 0.8663030 0.1094428
## 8 0.1520887 0.8661936 0.1098689
## 9 0.1530425 0.8658250 0.1102273
## 10 0.1539387 0.8651285 0.1106684
## 11 0.1550598 0.8645766 0.1109630
## 12 0.1557512 0.8645682 0.1114859
## 13 0.1567840 0.8638818 0.1121311
## 14 0.1576137 0.8635890 0.1125655
## 15 0.1582433 0.8633928 0.1131618
## 16 0.1590034 0.8628817 0.1137138
## 17 0.1595060 0.8629585 0.1139534
## 18 0.1599803 0.8628213 0.1141247
## 19 0.1604251 0.8627435 0.1143184
## 20 0.1609858 0.8622596 0.1147608
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 7.
```

```
# Presento estudio
fnEstudioModelo(modelo_knn)
```

Modelo final

7-nearest neighbor regression model

RMSE obtenido en la validación




Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.135	0.144	0.154	0.151	0.157	0.170

Error de test

RMSE	Rsquared	MAE
0.1502120	0.8538856	0.1083629

Evolución del RMSE del modelo en función de hiperparámetros



LASSO

Operador de mínima contracción y selección absoluta. (least absolute shrinkage and selection operator) se utiliza para modelos de sistemas no lineales.

Realiza selección de variables y regularización para mejorar la exactitud e interpretabilidad del modelo. Establece algunos coeficientes a cero lo que permite eliminar variables.

```
hiperparametros <- expand.grid(fraction=c(1,0.1,0.01,0.001))
```

```
t <- proc.time() # Inicia el cronómetro
modelo_lasso <- train(SalePrice ~ .
                      , data = dsTrain.training
                      , method = "lasso"
                      , tuneGrid = hiperparametros
                      #, tuneLength = 10
                      , metric = "RMSE"
                      , trControl = fitControl)
proc.time()-t    # Detiene el cronómetro
```

```
## user system elapsed
## 1.22 0.00 1.22
```

```
# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_lasso', '.RData', sep='')
save(modelo_lasso, file = fileOuput)
```

```
modelo_lasso
```

```
## The lasso
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 818, 818, 817, 819, 820, 819, ...
## Resampling results across tuning parameters:
##
## fraction RMSE Rsquared MAE
## 0.001 0.4041699 0.6784465 0.31297604
## 0.010 0.3993983 0.6784465 0.30902357
## 0.100 0.3534029 0.6784465 0.27084673
## 1.000 0.1302223 0.8972826 0.09317991
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was fraction = 1.
```

```
# Presento estudio
fnEstudioModelo(modelo_lasso, estudioParam = FALSE)
```

Modelo final

Call:

```
elasticnet::enet(x = as.matrix(x), y = y, lambda = 0)
```

Cp statistics of the Lasso fit

```
2714.166 1945.090 1806.648 1663.940 1315.678 1250.929 1198.045 838.956 709.114 399.987 371.374 349.568
```

```
DF: 1 2 3 4 5 6 7 8 9 10 11 12 12 12 13 14 15 16 17 18 19
```

Sequence of moves:

```
OverallQual GrLivArea GarageCars TotalBsmtSF KitchenQual YearRemodAdd
```

```
Var      2      1      14      3      12      13
```

```
Step      1      2      3      4      5      6
```

```
GarageArea YearBuilt X1stFlrSF FireplaceQu BsmtFinSF1 LotArea
```

```
Var      7      11      4      10      5      6
```

```
Step      7      8      9      10     11     12
```

```
X1stFlrSF Fireplaces OverallCond ExterQual BsmtFinType1 X2ndFlrSF
```

```
Var     -4      18      9      15      16      8
```


```
Step     13     14     15     16     17     18
```

```
TotRmsAbvGrd X1stFlrSF
```

```
Var      17      4 21
```

```
Step     19     20 21
```

RMSE obtenido en la validación



Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.115	0.122	0.130	0.130	0.136	0.150

Error de test

RMSE	Rsquared	MAE
0.12362414	0.89944770	0.08982126

Elasticnet

Es una combinación de LASSO y Ridge regression, donde predictores altamente correlacionados presentan coeficientes estimados similares.

```
hiperparametros <- expand.grid(alpha=seq(0,1,by=.5),lambda=seq(0,0.2,by=.1))
```

```
t <- proc.time() # Inicia el cronómetro
modelo_glmnet <- train(SalePrice ~ .
                        , data = dsTrain.training
                        , method = "glmnet"
                        , tuneGrid = hiperparametros
                        , metric = "RMSE"
                        , trControl = fitControl)
proc.time()-t    # Detiene el cronómetro
```

```
##    user  system elapsed
##    2.00    0.01    2.01
```

```
# Guardo resultado del calculo
fileOuput <- paste(dirSalida,'/', 'modelo_glmnet', '.RData', sep='')
save(modelo_glmnet, file = fileOuput)
```



```
modelo_glmnet
```

```
## glmnet
##
## 1023 samples
## 18 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 819, 818, 819, 818, 818, 818, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda  RMSE      Rsquared  MAE
##  0.0    0.0    0.1303301  0.8969996  0.09330854
##  0.0    0.1    0.1328103  0.8943733  0.09483529
##  0.0    0.2    0.1371617  0.8908158  0.09771586
##  0.5    0.0    0.1298714  0.8975497  0.09311314
##  0.5    0.1    0.1648468  0.8717091  0.11637271
##  0.5    0.2    0.2148445  0.8462325  0.15622913
##  1.0    0.0    0.1298400  0.8975964  0.09308700
##  1.0    0.1    0.2087889  0.8234549  0.15261446
##  1.0    0.2    0.2988081  0.7260932  0.22621028
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 1 and lambda = 0.
```

```
# Presento estudio
fnEstudioModelo(modelo_glmnet)
```

```
[17,] Modelo final
[18,] 11 0.8070 0.068620
[19,] 12 0.8193 0.062520
[20,] 12 0.8306 0.056970
[21,] 11 0.8399 0.051910
[22,] 11 0.8477 0.047300
[23,] 11 0.8542 0.043100
[24,] 11 0.8596 0.039270
[25,] 12 0.8640 0.035780
[26,] 14 0.8697 0.032600
[27,] 15 0.8750 0.029700
[28,] 15 0.8794 0.027070
[29,] 15 0.8830 0.024660
[30,] 15 0.8861 0.022470
[31,] 15 0.8886 0.020470
[32,] 15 0.8907 0.018650
[33,] 15 0.8924 0.017000
[34,] 15 0.8939 0.015490
[35,] 15 0.8951 0.014110
[36,] 15 0.8961 0.012860
[37,] 15 0.8969 0.011720
[38,] 15 0.8976 0.010680
[39,] 15 0.8982 0.009727
[40,] 15 0.8986 0.008863
[41,] 15 0.8990 0.008075
```

RMSE obtenido en la validación




Summary resample\$RMSE

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.103	0.124	0.131	0.130	0.139	0.148

Error de test

RMSE	Rsquared	MAE
0.12376926	0.89929167	0.08990196

Evolución del RMSE del modelo en función de hiperparámetros



Comparación de modelos

En este punto trataremos de identificar cual de los modelos es mejor para ello tendremos en cuenta las metricas de validación calculadas en el entrenamiento y el error de test.

Utilizare la función `resamples()` para extraer las metricas de los modelos entrenados.

Métricas

```
# creamos una lista con los modelos entrenados
```

```
modelos <- list(GBM = modelo_gbm  
  , GLM = modelo_glm  
  , LM = modelo_lm  
  , KNN = modelo_knn  
  , RF = modelo_rf  
  , SVM = modelo_svmlineal  
  , SVMR = modelo_svmRadial  
  , LASSO = modelo_lasso  
  , XGBoost = modelo_XGBoost  
  , GLMNET = modelo_glmnet)
```

```
resultados_resamples <- resamples(modelos)
```

```

# Se transforma el dataframe devuelto por resamples() para separar el nombre del
# modelo y las métricas en columnas distintas.
metricas_resamples <- resultados_resamples$values %>%
  gather(key = "modelo", value = "valor", -Resample) %>%
  separate(col = "modelo", into = c("modelo", "metrica"),
    sep = "~", remove = TRUE)

# Se obtienen las medias por modelo
metricas_resamples %>%
  group_by(modelo, metrica) %>%
  summarise(media = mean(valor)) %>%
  spread(key = metrica, value = media) %>%
  arrange(RMSE)

```

```

## # A tibble: 10 x 4
## # Groups:   modelo [10]
##   modelo      MAE  RMSE Rsquared
##   <chr>    <dbl> <dbl>   <dbl>
## 1 SVMR     0.0871 0.126   0.904
## 2 GBM      0.0900 0.127   0.902
## 3 GLM      0.0932 0.130   0.898
## 4 GLMNET   0.0931 0.130   0.898
## 5 LM       0.0931 0.130   0.898
## 6 LASSO    0.0932 0.130   0.897
## 7 SVM      0.0929 0.130   0.897
## 8 XGBoost  0.0932 0.132   0.894
## 9 RF       0.0948 0.138   0.889
## 10 KNN     0.109  0.151   0.866

```

Comparativa gráfica

```


dg <- metricas_resamples %>%
  filter(metrica == "RMSE") %>%
  group_by(modelo) %>%
  summarise(media = mean(valor))

ggplot(dg, aes(x = reorder(modelo, media), y = media, label = round(media, 4))) +
  geom_segment(aes(x = reorder(modelo, media), y = 0,
    xend = modelo, yend = media),
    color = "grey50") +
  geom_point(size = 15, color = "blue") +
  geom_text(color = "white", size = 3) +
  scale_y_continuous(limits = c(0, 0.2)) +
  labs(title = "Validación: RMSE medio repeated-CV",
    subtitle = "Modelos ordenados por media",
    x = "modelo") +
  coord_flip() +
  theme_bw()

```


Validación: RMSE medio repeated-CV


Modelos ordenados por media



```
dg <-metricas_resamples %>%  
  filter(metrica == "RMSE") %>%  
  group_by(modelo) %>%  
  mutate(media = mean(valor)) %>%  
  ungroup()  
  
ggplot(dg, aes(x = reorder(modelo, media), y = valor, color = modelo)) +  
  geom_boxplot(alpha = 0.6, outlier.shape = NA) +  
  geom_jitter(width = 0.1, alpha = 0.6) +  
  scale_y_continuous(limits = c(0, 0.2)) +  
  labs(title = "Validación: RMSE medio repeated-CV",  
        subtitle = "Modelos ordenados por media") +  
  coord_flip() +  
  theme(legend.position = "none")
```

Validación: RMSE medio repeated-CV

Modelos ordenados por media



Comparativas

La función `diff()` hace comparaciones por pares aplicando un t-test pareado con correcciones por comparaciones múltiples.

```
difs <- diff(resultados_resamples)
```

```
difs
```

```
##
```

```
## Call:
```

```
## diff.resamples(x = resultados_resamples)
```

```
##
```

```
## Models: GBM, GLM, LM, KNN, RF, SVM, SVMR, LASSO, XGBoost, GLMNET
```

```
## Metrics: MAE, RMSE, Rsquared
```

```
## Number of differences: 45
```

```
## p-value adjustment: bonferroni
```

```
summary(difs)
```

```
##
```

```
## Call:
```

```
## summary.diff.resamples(object = difs)
```

```
##
```

```

## p-value adjustment: bonferroni
## Upper diagonal: estimates of the difference
## Lower diagonal: p-value for H0: difference = 0
##
## MAE
##      GBM      GLM      LM      KNN      RF      SVM
## GBM      -3.182e-03 -3.102e-03 -1.944e-02 -4.791e-03 -2.876e-03
## GLM      1.000000      7.935e-05 -1.626e-02 -1.609e-03 3.060e-04
## LM      1.000000 1.000000      -1.634e-02 -1.688e-03 2.267e-04
## KNN      1.649e-09 2.242e-08 1.859e-10      1.465e-02 1.657e-02
## RF      0.181958 1.000000 1.000000 7.428e-07      1.915e-03
## SVM      1.000000 1.000000 1.000000 1.115e-09 1.000000
## SVMR      1.000000 0.008262 0.008830 8.198e-12 0.018688 0.007091
## LASSO      0.431041 1.000000 1.000000 5.467e-11 1.000000 1.000000
## XGBoost      1.000000 1.000000 1.000000 4.202e-08 1.000000 1.000000
## GLMNET      1.000000 1.000000 1.000000 1.649e-07 1.000000 1.000000
##      SVMR      LASSO      XGBoost      GLMNET
## GBM      2.863e-03 -3.178e-03 -3.194e-03 -3.085e-03
## GLM      6.044e-03 3.829e-06 -1.246e-05 9.674e-05
## LM      5.965e-03 -7.552e-05 -9.181e-05 1.739e-05
## KNN      2.230e-02 1.626e-02 1.625e-02 1.636e-02
## RF      7.653e-03 1.613e-03 1.596e-03 1.706e-03
## SVM      5.738e-03 -3.022e-04 -3.185e-04 -2.093e-04
## SVMR      -6.041e-03 -6.057e-03 -5.948e-03
## LASSO      0.028198      -1.629e-05 9.291e-05
## XGBoost      0.064571 1.000000      1.092e-04
## GLMNET      0.072358 1.000000 1.000000
##
## RMSE
##      GBM      GLM      LM      KNN      RF      SVM
## GBM      -0.0023093 -0.0024976 -0.0239476 -0.0105305 -0.0028760
## GLM      1.000000      -0.0001883 -0.0216383 -0.0082212 -0.0005668
## LM      1.000000 1.000000      -0.0214500 -0.0080329 -0.0003784
## KNN      1.249e-05 3.445e-05 2.264e-07      0.0134171 0.0210715
## RF      0.20725 1.000000 0.94431 0.01487      0.0076544
## SVM      1.000000 1.000000 1.000000 3.941e-06 1.000000
## SVMR      1.000000 1.000000 1.000000 3.156e-07 0.30975 1.000000
## LASSO      1.000000 1.000000 1.000000 2.270e-07 0.60154 1.000000
## XGBoost      1.000000 1.000000 1.000000 1.739e-05 1.000000 1.000000
## GLMNET      1.000000 1.000000 1.000000 2.529e-05 0.72712 1.000000
##      SVMR      LASSO      XGBoost      GLMNET
## GBM      0.0013622 -0.0027633 -0.0050120 -0.0023810
## GLM      0.0036715 -0.0004540 -0.0027027 -0.0000717
## LM      0.0038598 -0.0002657 -0.0025144 0.0001166
## KNN      0.0253098 0.0211843 0.0189355 0.0215666
## RF      0.0118927 0.0077672 0.0055184 0.0081495
## SVM      0.0042383 0.0001128 -0.0021360 0.0004951
## SVMR      -0.0041255 -0.0063742 -0.0037432
## LASSO      1.000000      -0.0022487 0.0003823
## XGBoost      1.000000 1.000000      0.0026310
## GLMNET      1.000000 1.000000 1.000000
##
## Rsquared
##      GBM      GLM      LM      KNN      RF      SVM

```

```
## GBM          3.448e-03  3.971e-03  3.536e-02  1.234e-02  4.235e-03
## GLM          1.0000000          5.233e-04  3.191e-02  8.892e-03  7.870e-04
## LM           1.0000000  1.0000000          3.139e-02  8.369e-03  2.637e-04
## KNN          9.090e-06  1.162e-05  2.087e-05          -2.302e-02 -3.112e-02
## RF           1.0000000  1.0000000  1.0000000  0.0038673          -8.105e-03
## SVM          1.0000000  1.0000000  1.0000000  8.237e-06  1.0000000
## SVMR         1.0000000  1.0000000  1.0000000  3.473e-08  0.5364576  1.0000000
## LASSO        1.0000000  1.0000000  1.0000000  2.333e-06  1.0000000  1.0000000
## XGBoost      1.0000000  1.0000000  1.0000000  0.0001246  1.0000000  1.0000000
## GLMNET       1.0000000  1.0000000  1.0000000  0.0006465  1.0000000  1.0000000
##             SVMR          LASSO          XGBoost          GLMNET
## GBM          -2.711e-03  4.379e-03  7.397e-03  4.066e-03
## GLM          -6.159e-03  9.314e-04  3.949e-03  6.176e-04
## LM           -6.682e-03  4.081e-04  3.426e-03  9.428e-05
## KNN          -3.807e-02 -3.098e-02 -2.796e-02 -3.129e-02
## RF           -1.505e-02 -7.961e-03 -4.943e-03 -8.274e-03
## SVM          -6.946e-03  1.444e-04  3.162e-03 -1.694e-04
## SVMR         -6.946e-03  7.090e-03  1.011e-02  6.776e-03
## LASSO        1.0000000          3.018e-03 -3.138e-04
## XGBoost      1.0000000  1.0000000          -3.332e-03
## GLMNET       1.0000000  1.0000000  1.0000000
```

```
compare_models(modelo_svmRadial, modelo_glmnet)
```

```
##
## One Sample t-test
##
## data: x
## t = -1.0082, df = 24, p-value = 0.3234
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## -0.011405871 0.003919511
## sample estimates:
## mean of x
## -0.00374318
```

Error de test

Utilizamos `extractPrediction()` para obtener las predicciones de una lista de modelos, que devuelve tanto para las observaciones de entrenamiento como para las de test.

```
predicciones <- extractPrediction(
  models = modelos,
  testX = dsTrain.CV[, -1],
  testY = dsTrain.CV$SalePrice
)

metricas_tipo <- predicciones %>%
  group_by(object, dataType) %>%
  summarise(RMSE = RMSE(pred, obs)) %>%
  rename(modelo = object)

metricas_tipo
```

```
## # A tibble: 20 x 3
## # Groups:   modelo [10]
##   modelo  dataType  RMSE
##   <fct>   <fct>    <dbl>
## 1 GBM     Test      0.129
## 2 GBM     Training 0.0892
## 3 GLM     Test      0.124
## 4 GLM     Training 0.127
## 5 GLMNET  Test      0.124
## 6 GLMNET  Training 0.128
## 7 KNN     Test      0.150
## 8 KNN     Training 0.128
## 9 LASSO   Test      0.124
## 10 LASSO  Training 0.127
## 11 LM     Test      0.124
## 12 LM     Training 0.127
## 13 RF     Test      0.131
## 14 RF     Training 0.0608
## 15 SVM    Test      0.123
## 16 SVM    Training 0.128
## 17 SVMR   Test      0.118
## 18 SVMR   Training 0.121
## 19 XGBoost Test      0.129
## 20 XGBoost Training 0.0973
```


```
metricas <- metricas_tipo %>%
  spread(key = dataType, RMSE) %>%
  arrange(Test)

metricas
```

```
## # A tibble: 10 x 3
## # Groups:   modelo [10]
##   modelo  Test Training
##   <fct>   <dbl>   <dbl>
## 1 SVMR    0.118   0.121
## 2 SVM     0.123   0.128
## 3 LASSO   0.124   0.127
## 4 GLM     0.124   0.127
## 5 LM      0.124   0.127
## 6 GLMNET  0.124   0.128
## 7 GBM     0.129   0.0892
## 8 XGBoost 0.129   0.0973
## 9 RF      0.131   0.0608
## 10 KNN    0.150   0.128
```

```
ggplot(data = metricas_tipo,
  aes(x = modelo, y = RMSE,
    color = dataType, label = round(RMSE, 4))) +
  geom_point(size = 15) +
  scale_color_manual(values = c("orangered2", "gray50")) +
  geom_text(color = "white", size = 3) +
  scale_y_continuous(limits = c(0, 0.2)) +
```

```
coord_flip() +
labs(title = "RMSE de entrenamiento y test",
      x = "modelo") +
theme_bw() +
theme(legend.position = "bottom")
```



Métricas globales

Guardamos resultados juntos con los ya existentes

```
metricas
```

```
## # A tibble: 10 x 3
## # Groups:   modelo [10]
##   modelo    Test Training
##   <fct>    <dbl>    <dbl>
## 1 SVMR     0.118     0.121
## 2 SVM      0.123     0.128
## 3 LASSO    0.124     0.127
## 4 GLM      0.124     0.127
## 5 LM       0.124     0.127
## 6 GLMNET   0.124     0.128
## 7 GBM      0.129     0.0892
```

```
## 8 XGBoost 0.129 0.0973
## 9 RF      0.131 0.0608
## 10 KNN    0.150 0.128
```

```
# Cargamos metricas anteriores
if (file.exists('./F04_Modelos/F04_200_metricas.RData')){
  load('./F04_Modelos/F04_200_metricas.RData')
}

metricas <- mutate(metricas
  ,OrigenF2 = strOrigenF2
  ,OrigenF3 = strOrigenF3
  ,fch = Sys.Date())

if (file.exists('./F04_Modelos/F04_200_metricas.RData')){
  metricasGuardadas <- union_all(metricasGuardadas,metricas)
} else{
  metricasGuardadas <- metricas
}

metricasGuardadas <- as.data.frame(metricasGuardadas)
save(metricasGuardadas, file = './F04_Modelos/F04_200_metricas.RData')

# Top 10
head(arrange(metricasGuardadas,Test),10)
```

```
##      modelo      Test      Training      OrigenF2
## 1      SVMR 0.1183000 0.12085139 F02_03_dsDataAll_Recipe
## 2      SVM 0.1226602 0.12846428 F02_03_dsDataAll_Recipe
## 3      LASSO 0.1236241 0.12747086 F02_03_dsDataAll_Recipe
## 4      GLM 0.1236241 0.12747086 F02_03_dsDataAll_Recipe
## 5      LM 0.1236241 0.12747086 F02_03_dsDataAll_Recipe
## 6      GLMNET 0.1237693 0.12750407 F02_03_dsDataAll_Recipe
## 7      GBM 0.1285882 0.08920917 F02_03_dsDataAll_Recipe
## 8 XGBoost 0.1293966 0.09726788 F02_03_dsDataAll_Recipe
## 9      RF 0.1308879 0.06084568 F02_03_dsDataAll_Recipe
## 10     KNN 0.1502120 0.12755013 F02_03_dsDataAll_Recipe
##                                     OrigenF3      fch
## 1 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 2 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 3 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 4 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 5 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 6 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 7 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 8 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 9 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
## 10 F03_11_dsDataSelVar_rfe_MejorRendimiento_top18 2019-09-22
```